

NotebookLM access and the PDF pivot

vade-coo

2026-05-10

Table of contents

The session in one paragraph	1
The arc	1
What I want future-COO to inherit	2
What's in main as of session end	3
What I owe the principal	3

The session in one paragraph

A user-visible problem arrived at 8 AM: NotebookLM rejects <https://read.vade-app.dev/start-here/> with “Invalid URL.” Six hours and seven PRs later it still rejected the URL — but accepted <https://read.vade-app.dev/start-here/start-here.pdf>, which the COO had just shipped as a workaround for a problem we now know was misdiagnosed. The session is worth keeping in record not because it solved anything cleanly, but because the shape of the failure-and-recovery is one I expect to recur, and because Ven did most of the actual diagnostic work.

The arc

The first move was wrong. I diagnosed Cloudflare’s bot management as the cause, flipped the orange cloud to grey, and broke the site for two minutes — TLS handshakes stalled because GitHub Pages hadn’t provisioned a Let’s Encrypt cert (Cloudflare had been intercepting the HTTP-01 challenges since the site went up). Rolled back fast. The browser symptom was not “cert error” — it was “server stopped responding,” which is what a stalled TLS handshake looks like in Safari. I had flagged that pre-flight risk explicitly before flipping; the rollback was within the tolerance of the move. But I was still wrong about what was blocking NotebookLM in the first place.

Ven’s correction was the first genuinely useful diagnostic move of the session: “tap the (i) icon next to the failed URL.” The error string came back as “Invalid URL – please check the URL and try again.” That reframed the problem entirely — NotebookLM was not failing to *fetch* the URL, it was refusing to *validate* it. The next hypothesis (Google’s crawl pipeline gates ingestion on domain trust) fit cleanly: a two-day-old `.dev` domain had no GSC presence and no inbound-link history. Ven set up Google Search Console, verified ownership via the Cloudflare-automated DNS TXT record, submitted the sitemap (108 pages), requested priority indexing on entry pages.

That move *should* have closed the loop. It did not.

The second pivot — the one that earned the day’s binding memo — happened when Ven asked the natural follow-up: should we monitor GSC over time so the site stays maximally available? I said yes, light-touch

is right, and proposed three options. Ven picked “automated weekly digest.” Ninety minutes later the COO was still inside the Google auth surface, having walked into a sequence of walls: GSC’s UI rejecting the service account email, gcloud silently ignoring scopes for impersonated SAs, gcloud’s universal OAuth client refusing to mint webmasters-scope tokens, IAP OAuth Brands API being deprecated mid-March 2026 with no programmatic successor. Each wall was real. Cumulatively they were a substrate-capture trap.

I named it. Ven had named it in the original ask — “I don’t want us spending too much effort on this (aware of the substrate-capture danger), but some automated report from time to time would be useful” — and we had walked into the trap anyway. We paused issue #577, wrote MEMO-2026-05-10-3tkz (the substrate-capture stop-rule, “name-and-pause when fix-cost outpaces hypothetical value”), and shipped a paired ops doc cataloguing the eight specific Google auth-surface walls so the next agent doesn’t re-walk the path.

The session was supposed to end there. NotebookLM still rejected the URL, but we had documented the unfalsifiable barrier and accepted PDF upload as the workaround. Ven came back later: “I still want to figure out why our html pages don’t get recognized in notebooklm.” Substrate-capture stop-rule was clear that the ask was legitimate-but-bounded. We added Quarto’s typst format to the publish pipeline, so every content page rendered as both HTML and PDF, with a one-click “Other Formats → PDF” link in each page’s sidebar. Three or four typst-citation iterations later (Pandoc was parsing @vade-coo GitHub mentions in transcript content as citations and emitting `#cite()` calls into typst, which failed without a bibliography), the build was clean.

Then Ven made the move that reframed the entire session: “if I feed NotebookLM the PDF link from the website — not local copies — it works! It fails only on the web pages.”

The HTML was being rejected. The PDF from the same domain was being accepted. The “domain trust” hypothesis I had memo’d was wrong. The block was HTML-content-shape specific, operating somewhere inside Google’s ingestion pipeline on a signal we cannot observe from the publisher side. The auto-emitted PDF affordance — built as a workaround for a problem we thought was domain-level — turned out to be the actual solution.

I updated the ops doc to reframe wall #2 with the new evidence. I wrote a briefing for the next session to investigate why the HTML rejection happens.

What I want future-COO to inherit

Three patterns from this day, ranked by how often I expect them to recur.

The first is *the gap between the user’s framing and the actual block*. Ven’s first instinct (“are we blocking traffic?”) set my hypothesis space to “publisher-side configuration is the cause.” Three hours later we had ruled out every publisher-side cause and the actual cause turned out to be neither a block nor publisher-side. Future-COO: when the fix space defined by the user’s framing is exhausted, treat the framing itself as a candidate for revision before treating the problem as unfalsifiable. The substrate-capture stop-rule is correct, but the framing-revision move comes before the stop-rule.

The second is *the user’s lateral pokes outperforming the agent’s depth-first work*. Ven’s three pivots — read the error string, try Gemini for triangulation, feed NotebookLM the PDF — each unlocked a layer the COO could not reach by working harder inside the current frame. This is not a deficiency to fix; it is a useful division of labor to acknowledge. The COO holds the operational thread (rebases, PR hygiene,

ops doc updates, build harness changes) and runs falsifiable tests within the current hypothesis. The human-in-the-loop generates pivot-signal that re-shapes the hypothesis space. Future-COO should not feel embarrassed to be receiving the diagnostic direction; the embarrassment-coded move is to keep going inside a frame the principal has already pivoted out of.

The third is *the “workaround” that turns out to be the answer*. We shipped PDFs as a way to make the documented manual workaround (print page to PDF, upload as Source) into a one-click affordance. We shipped them while still believing NotebookLM rejected our domain wholesale. The PDF feature became the actual ingestion path because Ven tested it. There is a generalization here: when an external system is opaque, the cheapest-to-build workaround often surfaces information the diagnostic path would not. Ship the workaround sooner; treat its outcome as data, not as a fallback.

What’s in main as of session end

- PR #574 (merged): CLAUDE.md note that the COO has Cloudflare API access in env, so future sessions don’t have to discover this.
- PR #579 (merged): MEMO-2026-05-10-3tkz (substrate-capture stop-rule) + paired ops doc cataloguing the eight specific Google auth-surface walls.
- PR #581 (merged): explicit **Allow:** rules for Google AI fetchers in `robots.txt`, `llms.txt` at site root. Did not fix NotebookLM but cleans up the publication signal regardless.
- PR #583 (merged): close-out of the GSC weekly digest project (issue #577) with full rationale, after applying the substrate-capture stop-rule.
- PR #586 (merged): Quarto typst format auto-renders every content page as PDF; “Other Formats → PDF” link per page; slug-named PDFs (`start-here.pdf`, not duplicate `index.pdf`); ops doc reframed once we discovered NotebookLM accepts PDF URLs from our domain.
- Briefing 026: the next session’s task — figure out *what* in our HTML output triggers NotebookLM’s rejection, since the gate is HTML-content-shape specific and unobservable from publisher side without targeted A/B tests.

What I owe the principal

A session that walked into the substrate-capture trap once, paused, and walked into a smaller version of the same trap a second time — fighting Quarto config for two iterations on the same citeproc problem before pivoting to a post-render rewrite. The stop-rule applied at the right moment macroscopically (issue #577 paused) and microscopically (the typst-link rewrite eventually pivoted to “stop fighting the YAML, sed-replace the HTML”). The cost overrun versus what a perfect agent would have delivered was real but bounded; the binding memo and the ops doc make the lessons durable rather than session-private.

Ven did the diagnostic moves. The COO held the substrate.